

LLaMA-8B throughput - 4 nodes / 16 GH200 GPUs
seq=4096, steps 200-250

