

TE vs FA3 throughput scaling with sequence length
LLaMA-8B, 4 nodes / 16 GH200 GPUs, steps 200-250

